

Collaborators: None

## Project Overview

- **Goal:** How can we find natural groups among subsidized housing based on their key features?
- **Dataset:** [Subsidized Housing - Six Metro Areas - 2017](#), 347 KB (1,945 lines)

## Data Processing

- **Loading:** The function `load_cleaned_data()` opens the CSV via `csv::Reader` and `serde::Deserialize`
  - Skips invalid rows (logging errors) and returns a `Vec<HousingProperty>` with fields `total_units: u32`, `subsidy_count: u32`, and `owner_type: String`
- **Cleaning/transformations:** Removed rows with missing entries (and referenced the CSV in SPSS to double-check), selected only the variables Latitude, Longitude, TotalUnits, ActiveSubs, and OwnerType (but finalized with the last three), and wrote out to “cleaned\_subsidized\_housing.csv” from “data\_cleaning.py”

## Code Structure

- **Modules**
  - “data.rs” loads and deserializes the cleaned CSV into structs
  - “clustering.rs” converts structs into a scaled 2D dataset, runs k-means (k = 4), denormalizes centroids, and prints summary statistics
  - “plot.rs” creates a 1600x1200 PNG scatter plot of TotalUnits and ActiveSubs, colored by cluster, with a legend and ‘x’ markers on each centroid
  - “utils.rs” provides a function used to match cluster IDs to colors consistently
  - “main.rs” runs all of the above where the project loads, clusters, and plots
- **Key functions & types**
  - `struct HousingProperty` (“data.rs”) → represents one housing record with fields `total_units`, `subsidy_count`, `owner_type`
  - `load_cleaned_data(path) (“”)` → inputs file path to cleaned CSV and outputs `Vec<HousingProperty>` or error; it opens a file, iterates, and pushes valid entries
  - `to_ndarray_with_scales(properties)` (“clustering.rs”) → inputs slice of `HousingProperty` and outputs `(Array2<f64>, min[0, 1], range[0, 1])`; it fills raw values, computes the minimum and maximum per column, and applies a normalization formula
  - `cluster_properties(properties, k) (“”)` → inputs data slice and cluster count and outputs `Vec<usize>` of labels; it calls the scaler, fits k-means, predicts labels,

and computes denormalized centroids, sizes, and owner-type distributions before printing

- *plot\_clusters(properties, labels)* (“plot.rs”) → inputs slice of properties and labels and outputs *()* or error; it automatically increments a new filename, builds *BitMapBackend*, draws points by cluster with legend entries, computes centroids and overlays markers, renders a legend, and saves
- *get\_cluster\_color(cluster\_id)* (“utils.rs”) → inputs cluster index and outputs *RGBColor* for plotting
- **Main workflow**
  - *main.rs* calls *load\_cleaned\_data* → *cluster\_properties* → *plot\_clusters*
  - Each stage uses only the public functions above, keeps data in memory, and communicates via return values and printed logs

## Tests

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\jiayi> cd "C:\Users\jiayi\section-8-rust"
PS C:\Users\jiayi\section-8-rust> cargo test
   Compiling section-8-rust v0.1.0 (C:\Users\jiayi\section-8-rust)
   Finished `test` profile [unoptimized + debuginfo] target(s) in 12.39s
   Running unittests src\main.rs (target\debug\deps\section_8_rust-e549bf54f62536b4.exe)

running 4 tests
test utils::tests::test_get_cluster_color_diff ... ok
test clustering::tests::test_two_clear_clusters ... ok
test data::tests::test_load_cleaned_data_simple ... ok
test plot::tests::test_plot_clusters_runs ... ok

test result: ok. 4 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out; finished in 1.15s

PS C:\Users\jiayi\section-8-rust> |
```

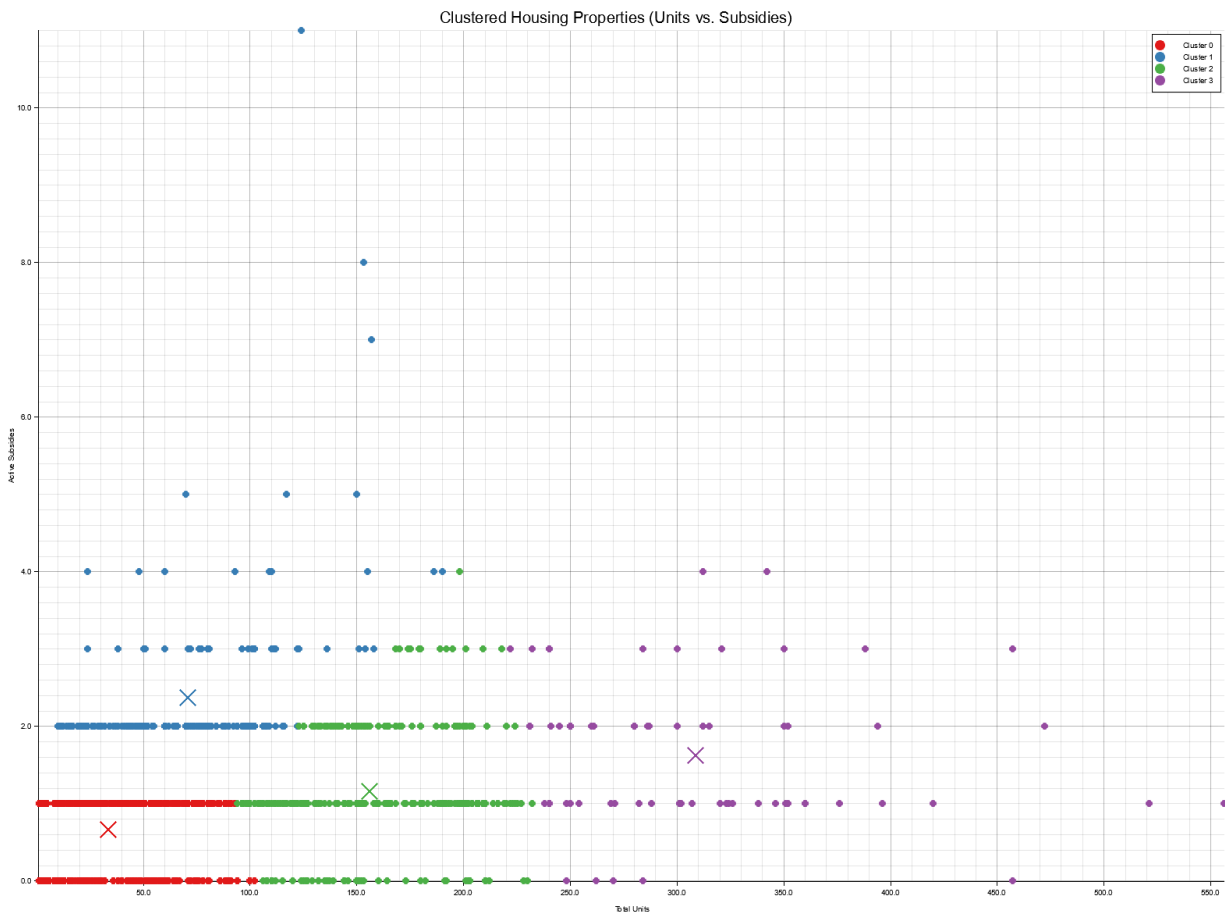
- For “**utils.rs**”: *test\_get\_cluster\_color\_diff()* checks whether different cluster IDs map to different colors on the RGB scale
- For “**clustering.rs**”: *test\_two\_clear\_clusters()* checks if *cluster\_properties* groups clearly distinct sample points into exactly two clusters
- For “**data.rs**”: *test\_load\_cleaned\_data\_simple()* checks that CSV parsing works on a basic example as a behavior scan
- For “**plot.rs**”: *test\_plot\_clusters\_runs()* checks for *plot\_clusters* execution doesn’t have errors and writes a PNG output file

# Results

```
Running 'target\release\section-8-rust.exe'
Loaded 1145 cleaned housing entries.
Cluster centroids:
Cluster 0:
Total Units: 34
Active Subsidies: 1
Cluster 1:
Total Units: 71
Active Subsidies: 2
Cluster 2:
Total Units: 156
Active Subsidies: 1
Cluster 3:
Total Units: 309
Active Subsidies: 2

Cluster sizes:
Cluster 0: 583 properties
Cluster 1: 200 properties
Cluster 2: 294 properties
Cluster 3: 68 properties

OwnerType distribution by cluster:
Cluster 0:
For Profit: 360
Multiple: 10
Non-Profit: 213
Cluster 1:
For Profit: 107
Multiple: 23
Non-Profit: 70
Cluster 2:
For Profit: 223
Multiple: 20
Non-Profit: 51
Cluster 3:
For Profit: 52
Multiple: 4
Non-Profit: 12
Saved plot to output/clusters_11.png
PS C:\Users\jiayi\section-8-rust> |
```



- Interpretation

- **Cluster 0** grouped a large number of small, least subsidized buildings where for-profits operate a majority and nonprofits fill in the rest; this suggests a widespread baseline assistance across both sectors
- **Cluster 1** grouped mid-sized developments with slightly more aid that are controlled mostly by for-profit companies; may suggest that gentrification is increasing
- **Cluster 2** grouped larger buildings receiving minimal aid where for-profits hold a three-quarters supermajority ownership; this suggests more housing scarcity and inequity over the distribution of resources
- **Cluster 3** grouped a small amount of the largest buildings with higher total subsidies, also controlled primarily by for-profits; may suggest housing displacement if the owners choose profit over people

## Usage Instructions

1. **Clone**
  - a. git clone <https://github.com/section-8-rust>
  - b. cd section-8-rust
2. **Build/Test**
  - a. cargo test
  - b. cargo build
3. **Run**
  - a. cargo run --release
4. **Runtime**
  - a. Data loading + clustering + plotting should finish in about 3-5 seconds

## AI-Assistance Disclosure and Other Citations

- I used ChatGPT to help me navigate file I/O when writing in Python (using Pandas) to clean my dataset. I am used to using absolute paths, but since the context of this assignment was to build a publicly accessible project, that was not viable. For this, I asked the following: “How do I use a relative path? In the grand scheme of things, I need my entire project to be accessible through GitHub (I want to do a push-pull request to update the work I've done today) so I don't think using absolute paths is

convenient.” It then produced this:

Keep a folder structure like this in your GitHub repo:

```
css                                                                    Copy Edit

your-project/
├─ data/
│   └─ Subsidized_Housing_-_Six_Metro_Areas_-_2017.csv
│       └─ cleaned_subsidized_housing.csv ← your output
├─ scripts/
│   └─ data_cleaning.py
├─ src/
│   └─ main.rs
├─ Cargo.toml
└─ requirements.txt
```

When you `cd your-project` in your shell (or open that folder in Sublime), all paths below are relative to that root.

- Among other things, ChatGPT helped me debug cases where my “NaN” filters were having errors and helped me track input and output locations (very useful since I started from C:\Users\jiayi and was not doing it in my dedicated DS 210 folder)
- Found [this blog post](#) to be particularly useful when thinking of how to visualize k-means clustering with my specific dataset, since my proposal was rejected for lack of graphability (and this was done in R, so no, I did not use any of their code)